

Machine Learning Kernels in Julia

trthatcher

May 3, 2015

Let $k : \mathbb{R}^d \times \mathbb{R}^d \mapsto \mathbb{R}$ be a kernel function. There are several classes that can be identified. This focusses on kernels of the form:

$$k(\mathbf{x}, \mathbf{y}) = \kappa(z(\mathbf{x}, \mathbf{y}))$$

where $z : \mathbb{R}^d \times \mathbb{R}^d \mapsto \mathbb{R}$ and $\kappa : \mathbb{R} \mapsto \mathbb{R}$.

For the first derivatives:

$$\frac{\partial \kappa}{\partial \mathbf{x}} = \kappa'(z) \left[\frac{\partial z}{\partial \mathbf{x}} \right] \quad \text{and} \quad \frac{\partial \kappa}{\partial \mathbf{y}} = \kappa'(z) \left[\frac{\partial z}{\partial \mathbf{y}} \right]$$

Then, the mixed second derivatives are:

$$\begin{aligned} \frac{\partial^2 \kappa}{\partial \mathbf{x} \partial \mathbf{y}} &= \frac{\partial}{\partial \mathbf{x}} \left(\left[\frac{\partial \kappa}{\partial \mathbf{y}} \right] \right) \\ &= \frac{\partial}{\partial \mathbf{x}} \left(\kappa'(z) \left[\frac{\partial z}{\partial \mathbf{y}} \right] \right) \\ &= \left[\frac{\partial z}{\partial \mathbf{y}} \right] \left[\frac{\partial \kappa'(z)}{\partial \mathbf{x}} \right]^\top + \kappa'(z) \frac{\partial}{\partial \mathbf{x}} \left(\left[\frac{\partial z}{\partial \mathbf{y}} \right] \right) \\ &= \kappa''(z) \left[\frac{\partial z}{\partial \mathbf{y}} \right] \left[\frac{\partial z}{\partial \mathbf{x}} \right]^\top + \kappa'(z) \left[\frac{\partial^2 z}{\partial \mathbf{x} \partial \mathbf{y}} \right] \end{aligned}$$

1 Scalar Product Kernel: $z(\mathbf{x}, \mathbf{y}) = \mathbf{x}^\top \mathbf{y}$

The Scalar Product kernel is of the form:

$$k(x, y) = \kappa(x^\top y)$$

where κ is a function unique to each Scalar Product kernel. For example, the the Sigmoid kernel instance of κ is defined to be:

$$\kappa_s(z; a, c) = \tanh(az + c)$$

Let X and Y be data matrices with rows as data points. Then the kernel matrix K :

$$X = \begin{bmatrix} \mathbf{x}_1^\top \\ \mathbf{x}_2^\top \\ \vdots \\ \mathbf{x}_n^\top \end{bmatrix} \text{ and } Y = \begin{bmatrix} \mathbf{y}_1^\top \\ \mathbf{y}_2^\top \\ \vdots \\ \mathbf{y}_m^\top \end{bmatrix} \implies K = \begin{bmatrix} \kappa(\mathbf{x}_1^\top \mathbf{y}_1) & \cdots & \kappa(\mathbf{x}_1^\top \mathbf{y}_m) \\ \vdots & \ddots & \vdots \\ \kappa(\mathbf{x}_n^\top \mathbf{y}_1) & \cdots & \kappa(\mathbf{x}_n^\top \mathbf{y}_m) \end{bmatrix}$$

Given a set of vectors $S = \{\mathbf{z}_1, \dots, \mathbf{z}_n\}$ the gramian is defined as $[\mathbf{G}]_{ij} = \mathbf{z}_i^\top \mathbf{z}_j$. Therefore, if we have two sets of vectors S_x and S_y with corresponding data matrices X and Y , then we can define the gramian between them as $\mathbf{G}(\mathbf{X}, \mathbf{Y}) = \mathbf{X}\mathbf{Y}^\top$. This is really the upper right-hand corner of the gramian of the set $S_x \cup S_y$.

Finally, we may define the kernel matrix as a function of the gramian:

$$[K]_{ij} = \kappa([G_{XY}]_{ij})$$

Generally, the best approach to take for implementation is:

1. Calculate G_{XY} for matrix X and Y using GEMM or SYRK using BLAS (exploit symmetry of kernels). The option to use rows or columns as data points is left as an option (using trans = 'N' assumes rows are observations and 'T' assumes columns). This is $O(n^3)$ but is heavily optimized to reduce the coefficient.
2. Scan through the elements of $\mathbf{G}(\mathbf{X}, \mathbf{Y})$ and apply the κ function to each element, overwrite G_{XY} in the process. This is $O(n^2)$.
3. Return K

2 Squared Distance Kernel

The Squared Distance kernel is of the form:

$$k(x, y) = \kappa(\|x - y\|^2)$$

where κ is a function unique to each kernel. For example, the the Gaussian kernel instance of κ is defined to be:

$$\kappa_\Phi(z; \sigma) = \exp\left(\frac{-z}{2\sigma^2}\right)$$

2.1 Kernel Matrix

Let \mathbf{X} and \mathbf{Y} be data matrices with rows as data points. Then the kernel matrix K :

$$\mathbf{X} = \begin{bmatrix} \mathbf{x}_1^\top \\ \mathbf{x}_2^\top \\ \vdots \\ \mathbf{x}_n^\top \end{bmatrix} \text{ and } \mathbf{Y} = \begin{bmatrix} \mathbf{y}_1^\top \\ \mathbf{y}_2^\top \\ \vdots \\ \mathbf{y}_m^\top \end{bmatrix} \implies \mathbf{K} = \begin{bmatrix} \kappa(\|\mathbf{x}_1 - \mathbf{y}_1\|^2) & \cdots & \kappa(\|\mathbf{x}_1 - \mathbf{y}_m\|^2) \\ \vdots & \ddots & \vdots \\ \kappa(\|\mathbf{x}_n - \mathbf{y}_1\|^2) & \cdots & \kappa(\|\mathbf{x}_n - \mathbf{y}_m\|^2) \end{bmatrix}$$

For a vector \mathbf{x} and \mathbf{y} , define the lag (alternatively, the error) between the two vectors to be:

$$\text{lag}(x, y) = x - y = \epsilon$$

Then we may define kernels of this form as:

$$k(x, y) = \kappa(\text{lag}(\mathbf{x}, \mathbf{y}) \cdot \text{lag}(\mathbf{x}, \mathbf{y})) = \kappa(\epsilon^\top \epsilon) = \kappa(\|\mathbf{x} - \mathbf{y}\|^2)$$

Given a set of vectors $S = \{\mathbf{z}_1, \dots, \mathbf{z}_n\}$, the lagged gramian is defined as:

$$[\mathbf{G}_l]_{ij} = \|\mathbf{z}_i - \mathbf{z}_j\|^2$$

Similarly, for two data matrices \mathbf{X} and \mathbf{Y} :

$$[\mathbf{G}_l(\mathbf{X}, \mathbf{Y})]_{ij} = \|\mathbf{x}_i - \mathbf{y}_j\|^2$$

Where \mathbf{x}_i and \mathbf{y}_j are the data vectors in \mathbf{X} and \mathbf{Y} . The lagged gramian can be defined in terms of $\mathbf{G}(\mathbf{X}, \mathbf{Y})$, \mathbf{X} and \mathbf{Y} .

$$[\mathbf{G}_l(\mathbf{X}, \mathbf{Y})]_{ij} = [\text{diag}(\mathbf{G}(\mathbf{X}))]_i - 2[\mathbf{G}]_{ij} + [\text{diag}(\mathbf{G}(\mathbf{Y}))]_j$$

The function $\text{diag}(\mathbf{G}(\mathbf{X}))$ returns a vector of diagonal entries of \mathbf{G} . This is simply a vector where entry i is $\mathbf{x}_i^\top \mathbf{x}_i$.

2.2 Kernel Matrix Calculation - Implementation

Therefore, to compute the kernel matrix for a squared distance kernel:

1. Calculate $\mathbf{G}(\mathbf{X}, \mathbf{Y})$ for matrix \mathbf{X} and \mathbf{Y} using GEMM or SYRK using BLAS (exploit symmetry of kernels). This is $O(n^3)$.
2. Transform $\mathbf{G}(\mathbf{X}, \mathbf{Y})$ to $\mathbf{G}_l(\mathbf{X}, \mathbf{Y})$. This is $O(n^2)$.
3. Scan through the elements of $\mathbf{G}_l(\mathbf{X}, \mathbf{Y})$ and apply the κ function to each element, overwrite $\mathbf{G}_l(\mathbf{X}, \mathbf{Y})$ in the process. This is $O(n^2)$.
4. Return \mathbf{K}

2.3 Kernel Derivative

$$\begin{aligned} \frac{\partial z}{\partial \mathbf{x}} &= \frac{\partial \|\mathbf{x} - \mathbf{y}\|^2}{\partial \mathbf{x}} = \frac{\partial}{\partial \mathbf{x}} (\mathbf{x}^\top \mathbf{x} - 2\mathbf{x}^\top \mathbf{y} + \mathbf{y}^\top \mathbf{y}) = 2(\mathbf{x} - \mathbf{y}) \\ \frac{\partial z}{\partial \mathbf{y}} &= \frac{\partial \|\mathbf{x} - \mathbf{y}\|^2}{\partial \mathbf{y}} = \frac{\partial}{\partial \mathbf{y}} (\mathbf{x}^\top \mathbf{x} - 2\mathbf{x}^\top \mathbf{y} + \mathbf{y}^\top \mathbf{y}) = 2(\mathbf{y} - \mathbf{x}) \end{aligned}$$

Using the identity:

$$\left[\frac{\partial \mathbf{u}}{\partial \mathbf{v}} \right]_{ij} = \frac{\partial u_i}{\partial v_j}$$

The second mixed derivative is:

$$\frac{\partial^2 z}{\partial \mathbf{x} \partial \mathbf{y}} = \frac{\partial \|\mathbf{x} - \mathbf{y}\|^2}{\partial \mathbf{x} \partial \mathbf{y}} = \frac{\partial}{\partial \mathbf{x}} 2(\mathbf{y} - \mathbf{x}) = -2\mathbf{I}_d$$

Substituting in the above into the formula for kernel derivative:

$$\begin{aligned}\frac{\partial \kappa^2}{\partial \mathbf{x} \partial \mathbf{y}} &= \kappa''(z) \left[\frac{\partial z}{\partial \mathbf{y}} \right] \left[\frac{\partial z}{\partial \mathbf{x}} \right]^\top + \kappa'(z) \left[\frac{\partial z^2}{\partial \mathbf{x} \partial \mathbf{y}} \right] \\ &= -4\kappa''(z)(\mathbf{x} - \mathbf{y})(\mathbf{x} - \mathbf{y})^\top - 2\kappa'(z)\mathbf{I}_d\end{aligned}$$

Define $\boldsymbol{\epsilon} = \mathbf{x} - \mathbf{y}$, then this gives an element-wise formula of:

$$\left[\frac{\partial \kappa^2}{\partial \mathbf{x} \partial \mathbf{y}} \right]_{ij} = \begin{cases} -4\kappa''(\boldsymbol{\epsilon}^\top \boldsymbol{\epsilon}) \epsilon_i \epsilon_j & \text{if } i \neq j \\ -4\kappa''(\boldsymbol{\epsilon}^\top \boldsymbol{\epsilon}) \epsilon_i^2 - 2\kappa'(\boldsymbol{\epsilon}^\top \boldsymbol{\epsilon}) & \text{if } i = j \end{cases}$$

2.4 Kernel Derivative - Implementation

To compute $\left[\frac{\partial \kappa^2}{\partial \mathbf{x} \partial \mathbf{y}} \right]_{ij}$:

1. Compute array \mathbf{A} of element-wise differences for all coordinates in arrays \mathbf{X} and \mathbf{Y} .
2. Compute the lagged gramian matrix
3. Scan through the elements of \mathbf{A} and apply the κ functions
4. Return K